


Prisoner's Dilemma


Stefanie Bodnar
Evian Kassab
Tyler Johnson
Zach Shaver



Source: James Wilson

Introduction

- Problem Introduction
- Strategies & Implementation
- Hill-Climbing
- Genetic Algorithm
- Conclusion



Problem Introduction

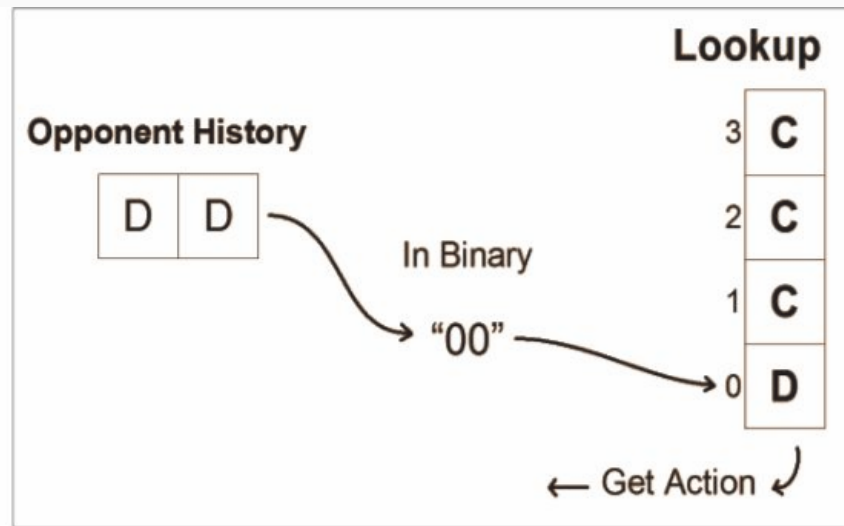
- Game & Real World Model
- Three Components
 - Player
 - Actions
 - Payoff Function
- Rationality & Irrationality
- Objective
 - Performance of Search Algorithm
 - Testing
 - Multiple Strategies
 - Compare & Contrast

		Bob	
		cooperate	defect
Ann	cooperate	-1, -1	-20, 0
	defect	0, -20	-10, -10



Strategies & Implementation


- Technology Used
 - IPDX Java Library
 - Provides basic programming interface for playing the game
- History of Moves
 - Stored as an integer, to be used as index in lookup table.
- Array Lookups
 - 1D - Opponent History
 - 2D - Opponent & Player History
- Evaluation Function
 - Generates heuristic by playing a set number of games against a training set and returning the average score



Strategies & Implementation

- Iterated Prisoner's Dilemma
 - Memory of Previous Moves
 - Goal: Best Lookup
 - Train algorithms to compete against specific, and general strategies.
 - Payoff Matrix
- Data Collection
 - Storing results in CSV to visualize with charts

	B: Cooperate	B: Defect
A: Cooperate	3,3	0,5
A: Defect	5,0	1,1



Hill-Climbing - Implementation

Input: Max Restarts, Max Sideways
Output : Lookup that is a Local Maximum

Lookup = *new Random Lookup.*

```
Repeat: {  
    Neighbour = getBestNeighbour().  
    if ( Neighbour's Score > Lookup's Score )  
        Store Neighbour as Lookup  
    else if ( Neighbour's Score < Lookup's Score OR Out of Restarts)  
        return Lookup  
    else  
        if ( Sideways Moves Available )  
            Do a Sideways Move  
        else  
            Reset Sideways Moves  
            Do a Restart  
}
```

- Operator Function - *getBestNeighbour()*
 - Gets the best neighbour by:
 - Flipping each action in current lookup and returning the neighbour with the best score.

- A local optimization approach
 - Start with a random lookup,
 - Search local neighbourhood for higher scoring tables,
 - Replace lookup with the best neighbour,
 - Repeat until no more are found.

Original Table:	CCCC (Score : 2.51)
Neighbour 1:	CCCD (Score : 2.48)
Neighbour 2:	CCDC (Score : 2.33)
Neighbour 3:	CDCC (Score : 2.65) <= Best Neighbour
Neighbour 4:	DCCC (Score : 2.56)


Hill Climbing - Findings

- 2D Lookup produces better results
 - Longer hill climbs, shorter histories & better solutions on average than 1D
- Solutions vary by length of history when trained against a variety of strategies



Hill Climbing - Findings

- Reaches a local max fairly quickly, which may not be optimal.
 - Solved by random restarts and sideways moves



Genetic Algorithm - Implementation

- Emulates evolution by selecting stronger individuals in a population for breeding.


Input: Population Size, Number of Generations, Number of Children per Couple

Output : Fittest Individual after Evolution

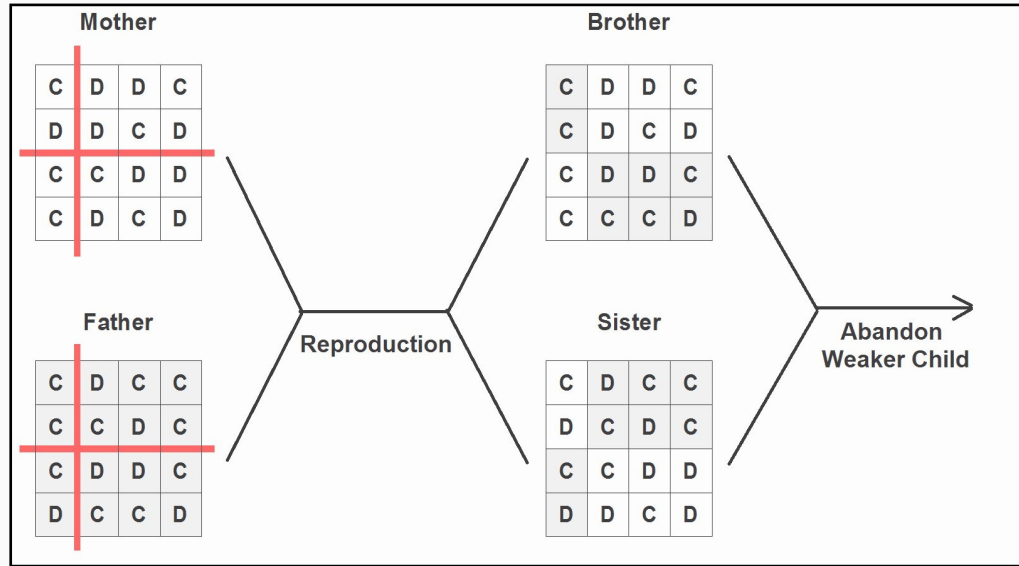
Population = new Population of Random Lookups

```
For: ( Each Generation ) {  
    NextGeneration = Empty Population.  
    For: ( Population size ) {  
        Father = selectRandomIndividual().  
        Mother = selectRandomIndividual().  
        Child = reproduce(Mother, Father).  
        Mutate Child with a small probability.  
        Add the Child to NextGeneration.  
    }  
    Population = NextGeneration  
}  
Return the Fittest Individual in Population
```

- selectRandomIndividual()*
 - Sorted by fitness, picks individual randomly according to a skewed distribution.




Genetic Algorithm - Implementation





- *reproduce*(Mother,Father)
 - After Mother and Father selected,
 - Splits them at random indices,
 - Exchanges their parts to form a child,
 - And adds “fitter” child to Next Generation
 - *Our algorithm also has a second mode that keeps both children.*



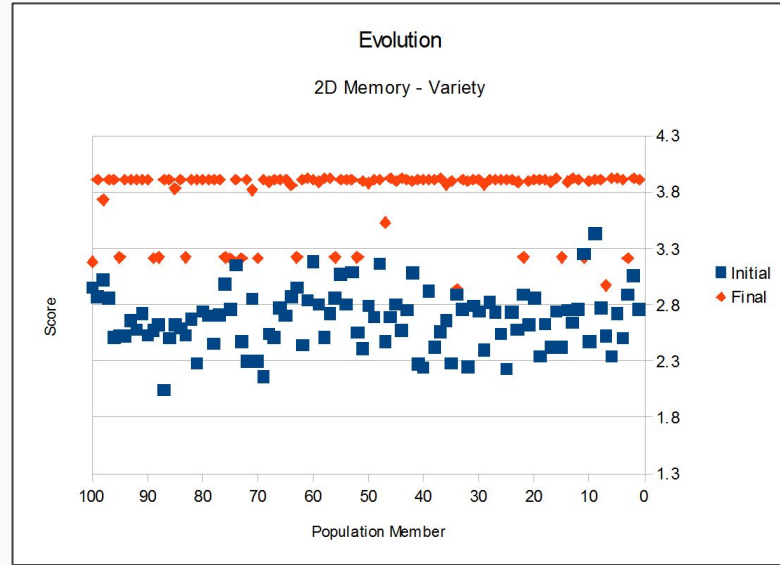

Genetic Algorithm - Findings



- The range of individual fitness declines over generations. Everyone gets fitter.
- Population's fitness peaks quickly
- 2D lookups perform better than 1D again.



Genetic Algorithm - Findings




- When facing:
 - Single strategy: There is a peak score, the optimal lookup.
 - Multiple strategies: Often no optimal lookup, but close!
- Entire population tends toward the exact same lookup table, which is the optimal solution.



Genetic Algorithm - Findings



- Changing Population Size & Number of Generations
 - No significant effect on fitness past 10 generations or 10 individuals.
- Changing Mutation Rate
 - More diverse population



Conclusion

Compare & Contrast


- Both algorithms prefer similar history depths, according to their lookup dimension.



Conclusion

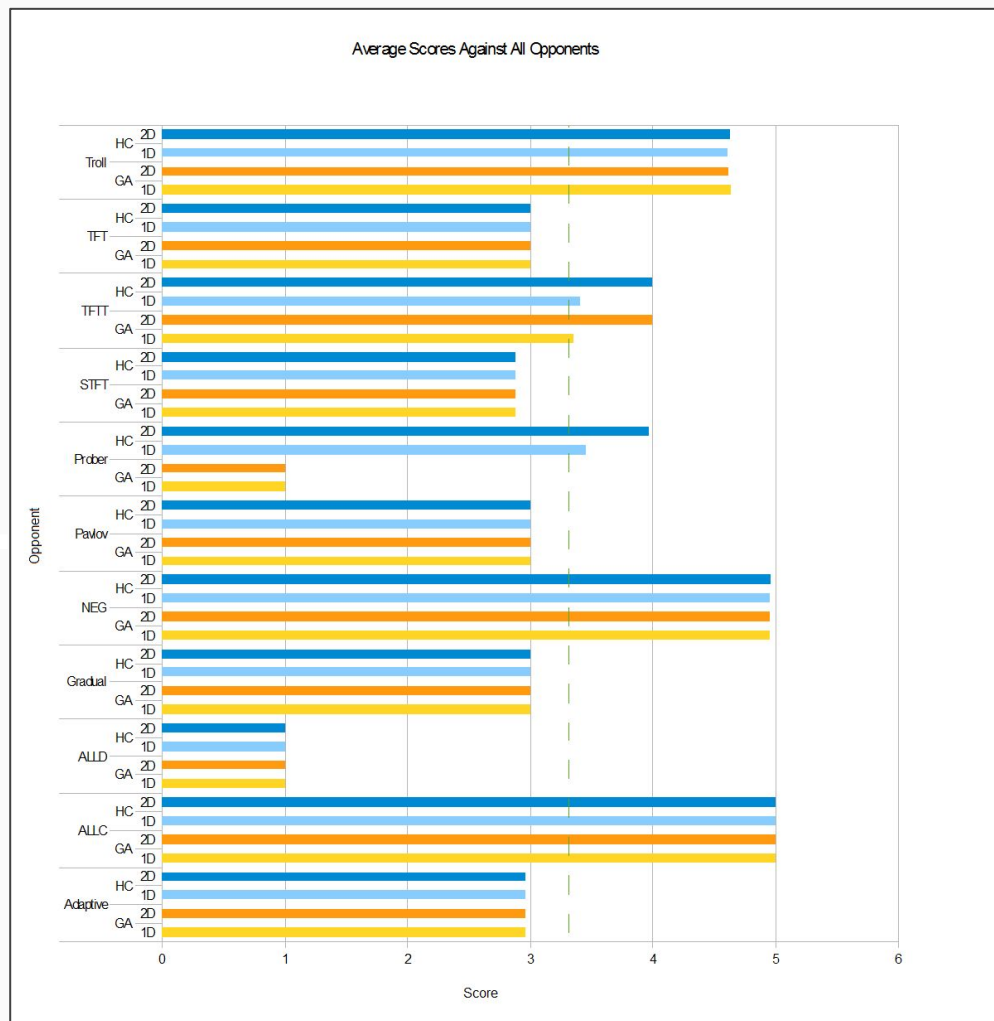
Compare & Contrast

- 2D Lookups based on opponent & player history produce better results, for both algorithms.



Conclusion

- Both algorithms consistently produced optimal lookup tables when trained against particular strategies.
- Algorithms were also trained against human-design strategies successfully.



Conclusion

- Both algorithms perform very well on the Prisoner's Dilemma problem.
 - Hill Climb can produce optimal tables quickly, while genetic algorithms will produce them more consistently.
- 2D array lookups, which look at opponent & player history together, often perform better than 1D lookups.

The End... *Any Questions?*

